


Step 24 — object interaction ($\kappa=0.8$): capacity wins in-distribution, equivariance wins across the (collapsed) group

(a) Interpolation: the MLP fits best
(VN cross-product cap; message gap $\times 1.36$)



(b) Extrapolation: only the prior is flat
(whole scene rotated off the training wedge)

