

Certified World Models: Predictability Across Configuration, Horizon, and Resolution

Scale buys interpolation; structure buys a certificate.

Hongbo Wang

Department of Mathematics, Stony Brook University, Stony Brook, NY 11794, USA

Abstract

Scaling a world model lowers its average error on the data manifold but issues **no guarantee about any specific unseen situation**. We give a different *kind* of result for **equivariant** latent world models — a **predictability certificate**: a provable, computable region of situations the model is guaranteed to handle *without further training or data*. The certificate spans three orthogonal axes at once: **configuration** (the exponentially large monoid $\langle S \rangle$ generated by k primitive symmetries, certified by checking the k generators), **horizon** T (how far ahead the guarantee reaches), and **resolution** ϵ (at what level of detail). A master theorem makes this precise: under exact equivariance with an orthogonal latent representation the whole-pipeline error is *invariant* over all of $\langle S \rangle$ (Theorem A), and a spectral refinement (Theorem B) shows that each latent channel is certified to a horizon $T_j(\epsilon) \sim \log(1/\epsilon)/\lambda_j$ set by its Lyapunov exponent — so the certified region is the **coarse-invariant, slow, low-composition** corner. We confirm all three axes at small scale: a model trained on the 6 generators of $\mathbb{Z}_2^6$ certifies all $2^6=64$ compositions to machine precision while a matched non-equivariant baseline degrades monotonically; a learned predictor recovers a chaotic Lyapunov exponent to within 0.4% and its certified horizon obeys the $\log(1/\epsilon)$ law; on an $\text{SO}(2)$ mechanical system the conserved (slowest) quantities organize into the architecturally invariant channels — a measured **Noether hinge** linking the configuration axis to the horizon axis; and a non-equivariant network scaled across an $88\times$ parameter range buys in-distribution interpolation (even beating the equivariant model there) yet over the unseen orbit stays $10\text{--}155\times$ above the equivariant floor and never reaches it. Finally, on **real PushT contact dynamics** — physics simulated by an engine we did not author — a learned equivariant world model’s multi-step rollout error is *exactly* flat over the orbit and competitive in-distribution, while no non-equivariant baseline across a $160\times$ parameter ladder reaches its out-of-distribution floor. *Scale buys interpolation; structure buys a certificate*. The result is a single criterion for *what* an equivariant world model can certifiably predict, and a structural reading of *why* celestial mechanics is forecastable for millennia while weather is not.

1. Introduction

Two debates organize modern world-model research. The first is *generation vs. abstraction*: pixel-space simulators (diffusion and video models) against latent predictive models that predict representations rather than pixels. The second is *scale vs. structure*: the view that brute-force scaling eventually wins, against the geometric-learning view that symmetry priors buy sample efficiency. Both are usually run as **degree** contests — who needs fewer pixels, who needs less data.

We argue the more useful question is a **kind** contest. Scaling produces a model with low *average* error on the data distribution, but it cannot certify anything about a *named, specific* situation the data did not cover. Structure can. For an equivariant world model we can write down, in advance and without further data, the exact set of situations on which the model’s behaviour is **guaranteed** — a *predictability certificate*. This reframes “is structure worth it?” from a benchmark race into a question about *what kind of statement each approach can make*.

The certificate has three orthogonal axes:

- **Configuration** $w \in \langle S \rangle$: the model generalizes across the exponentially large monoid generated by k primitive symmetries $S = \{g_1, \dots, g_k\}$, and we *certify the whole monoid by checking only the k generators*.
- **Horizon** T : how far into the future the guarantee reaches.
- **Resolution** ϵ : at what level of detail the guarantee holds.

Our contributions are:

1. **A three-axis master theorem** (§3): composition closure (k checks $\Rightarrow$ an exponential set), an exact certificate under exact equivariance (Theorem A) together with its **converse** (Lemma 2: the certificate is *equivalent* to equivariance, so no non-equivariant model has it at any size), and a spectral degradation law $T_j(\epsilon) \sim \log(1/\epsilon)/\lambda_j$ (Theorem B) whose certified region is the coarse-invariant-slow-low- $|w|$ corner — with a **quantitative scale-vs-structure separation** (§3.3: structure certifies the whole ϵ -independent orbit; the best L -Lipschitz non-equivariant learner certifies only an ϵ/L -tube around its data).
2. **The Noether hinge** (§4): the conjecture — *measured and confirmed on controlled systems* — that the group-invariant channels are the dynamically slow (long-horizon-certifiable) ones, linking the configuration axis to the horizon axis.
3. **Empirical confirmation of all three axes at small scale** (§5), including the headline contrast *scale buys interpolation; structure buys a certificate*, and a keystone validation on **real physics-engine contact dynamics** (PushT) where the certificate holds for a *learned* model of dynamics we did not design.

We are explicit about scope (§7): this is a mechanism-and-theory paper with 1–2-GPU proof-of-principle, not a scaled benchmark, and the hinge’s lift to *approximate* symmetry is open.

2. Setup

An encoder $E : \mathcal{X} \rightarrow \mathcal{Z}$ maps states to latents and an action-conditioned predictor $f : \mathcal{Z} \times \mathcal{A} \rightarrow \mathcal{Z}$ advances them; the true environment transition is $\Phi : \mathcal{X} \times \mathcal{A} \rightarrow \mathcal{X}$. A finite generating set $S = \{g_1, \dots, g_k\}$ generates a monoid $\langle S \rangle$ whose elements — words $w = g_{i_1} \dots g_{i_m}$ of length $|w| = m$ — act on states by $g \cdot x$, on latents by a representation $\rho(g)$, and on actions by $\sigma(g)$. Write $f^{(T)}(z; \vec{a})$ for the T -step latent rollout under actions $\vec{a} = (a_1, \dots, a_T)$, x_T^{true} for the T -step true state, and

$$\text{Err}_T(x; \vec{a}) = \|f^{(T)}(E(x); \vec{a}) - E(x_T^{\text{true}})\|$$

for the whole-pipeline prediction error, where $\|\cdot\|$ is the Euclidean norm on $\mathcal{Z}$ (the experiments use relMSE, its square).

We use the following assumptions, **each checkable on the k generators**:

- **(A1) Encoder equivariance:** $E(g_i \cdot x) = \rho(g_i) E(x)$.
- **(A2) Predictor equivariance:** $f(\rho(g_i)z, \sigma(g_i)a) = \rho(g_i)f(z, a)$.
- **(A3) The group is a symmetry of the dynamics:** $\Phi(g_i \cdot x, \sigma(g_i)a) = g_i \cdot \Phi(x, a)$, so that rolling out the *transformed* state equals transforming the *rolled-out* state, $(g_i \cdot x)_T^{\text{true}} = g_i \cdot x_T^{\text{true}}$ under $\sigma(g_i)\vec{a}$.
- **(A4) $\rho(g_i)$ orthogonal:** $\rho(g_i)^\top \rho(g_i) = I$ (norm-preserving on $\mathcal{Z}$).
- **(A5) Equivariant planner and invariant cost (closed-loop clause only):** the planner satisfies $\pi(\rho(g_i)z, \rho(g_i)z^*) = \sigma(g_i) \pi(z, z^*)$ and the cost is ρ -invariant, $J(\rho(g_i)z; \rho(g_i)z^*) = J(z; z^*)$. This is needed *only* for the closed-loop J identity, not for the prediction-error identity, and §5.6 instantiates exactly such a planner.

3. The Predictability Certificate

3.1 The configuration axis (Theorem A)

Lemma 1 (composition closure). If (A1)–(A4) hold on each generator $g_i \in S$, they hold on every word $w \in \langle S \rangle$, and ρ restricted to $\langle S \rangle$ is an orthogonal-matrix *monoid homomorphism*.

Proof. Define $\rho(w) := \rho(g_{i_1}) \cdots \rho(g_{i_m})$. For (A2), induct on m : $f(\rho(w)z, \sigma(w)a) = f(\rho(g_{i_1})[\rho(g_{i_2} \cdots)z], \sigma(g_{i_1})[\sigma(g_{i_2} \cdots)a]) = \rho(g_{i_1})f(\rho(g_{i_2} \cdots)z, \sigma(g_{i_2} \cdots)a) = \rho(w)f(z, a)$, applying the single-generator (A2) once per step. (A1) and (A3) lift identically by composition; (A4) lifts because a product of orthogonal matrices is orthogonal. $\square$

Hence k generator checks certify the entire (exponentially large) generated monoid.

Theorem A (exact configuration certificate). Under (A1)–(A4), for every $w \in \langle S \rangle$, every horizon T , and every action sequence $\vec{a}$,

$$\text{Err}_T(w \cdot x; \sigma(w)\vec{a}) = \text{Err}_T(x; \vec{a}), \quad \text{and, for an equivariant planner, } J(w \cdot x; w \cdot \text{goal}) = J(x; \text{goal}).$$

Proof. **(i) Rollout equivariance.** By induction on T : $f^{(1)} = f$ is equivariant by (A2) and Lemma 1; assuming $f^{(T)}(\rho(w)z; \sigma(w)\vec{a}) = \rho(w)f^{(T)}(z; \vec{a})$,

$$f^{(T+1)}(\rho(w)z; \sigma(w)\vec{a}) = f(f^{(T)}(\rho(w)z; \sigma(w)\vec{a}), \sigma(w)a_{T+1}) = f(\rho(w)f^{(T)}(z; \vec{a}), \sigma(w)a_{T+1}) = \rho(w)f^{(T+1)}(z; \vec{a}).$$

(ii) Predicted latent at $w \cdot x$. $f^{(T)}(E(w \cdot x); \sigma(w)\vec{a}) \stackrel{(A1)}{=} f^{(T)}(\rho(w)E(x); \sigma(w)\vec{a}) \stackrel{(i)}{=} \rho(w)f^{(T)}(E(x); \vec{a})$. **(iii) Target at $w \cdot x$.** $E((w \cdot x)_T^{\text{true}}) \stackrel{(A3)}{=} E(w \cdot x_T^{\text{true}}) \stackrel{(A1)}{=} \rho(w)E(x_T^{\text{true}})$. **(iv) Cancellation.** Subtracting (iii) from (ii) and using (A4) (an orthogonal ρ preserves the Euclidean norm),

$$\text{Err}_T(w \cdot x; \sigma(w)\vec{a}) = \|\rho(w)[f^{(T)}(E(x); \vec{a}) - E(x_T^{\text{true}})]\| = \|f^{(T)}(E(x); \vec{a}) - E(x_T^{\text{true}})\| = \text{Err}_T(x; \vec{a}).$$

The closed-loop identity is the same computation under (A5): the equivariant planner (lifted to words by Lemma 1) selects the $\sigma(w)$ -image of the same actions, and the ρ -invariant cost is unchanged. $\square$

Theorem A requires **(A3)** — the group must be a symmetry of the *dynamics*, not merely of the encoder. Where the dynamics break the symmetry, the certificate degrades by the amount of that break (Theorem B). The error is held *constant across the orbit*, not made small: $\text{Err}_T(x)$ itself can be large, and we report absolute errors alongside ratios throughout (§7).

Theorem A is sufficient; the next lemma shows equivariance is also *necessary* for the guarantee, so the certificate is not merely a consequence of equivariance but is **equivalent** to it — which is what makes the impossibility for non-equivariant models a theorem rather than a slogan.

Lemma 2 (the certificate characterizes equivariance). Let $\rho : G \rightarrow O(\mathcal{Z})$ act *freely* on an open $U \subseteq \mathcal{Z}$, and let $\mathcal{D}_G$ be the equivariant maps $\mathcal{Z} \rightarrow \mathcal{Z}$. If a predictor f 's one-step error $\|f - \Phi\|$ is orbit-constant on U for **every** target $\Phi \in \mathcal{D}_G$, then f is equivariant on U . *Proof.* Fix $z \in U$, g with $\rho(g)z \in U$. For any $c \in \mathcal{Z}$, freeness makes $\Phi(\rho(h)z) := \rho(h)c$ well-defined on the orbit of z (extend by 0 elsewhere; both pieces lie in $\mathcal{D}_G$ as $\rho(g)0 = 0$). Orbit-constancy gives $\|f(z) - c\| = \|f(\rho(g)z) - \rho(g)c\| = \|\rho(g)^{-1}f(\rho(g)z) - c\|$ (orthogonality of $\rho(g)$) for **all** c ; two points equidistant to every c coincide, so $f(\rho(g)z) = \rho(g)f(z)$. $\square$ (For G compact with closed orbits the interpolant can be taken continuous, so the statement holds against continuous dynamics, not only the full algebraic class.) With Theorem A this gives a **characterization**: orbit-constant error against every equivariant target $\iff f$ equivariant. Hence **no non-equivariant model possesses the configuration certificate, at any size** — the architectural impossibility invoked in §6–§7 is a theorem. The result is elementary (one line of Hilbert-space geometry once freeness frees the probe c); its role is to *pin down* what the certificate is, not to add machinery.

3.2 The horizon and resolution axes (Theorem B)

Theorem B (spectral degradation). Relax exactness: suppose the per-generator encoder residual is $\epsilon_{\max} = \max_i \sup_x \|E(g_i \cdot x) - \rho(g_i)E(x)\|$, the per-step predictor error is δ , and the latent map's Jacobian is, on the ℓ -isotypic channels j , locally diagonalized with multipliers e^{λ_j} (Lyapunov exponents). A Grönwall/Lipschitz accumulation — the word error enters m times (Lemma 1 composes m approximate generators) and the per-step error compounds over T steps amplified by the channel multiplier — gives, for constants c_j depending on the local geometry,

$$|\text{Err}_T(w \cdot x) - \text{Err}_T(x)| \leq \sum_{\text{channels } j} c_j (m \epsilon_{\max} + T \delta) e^{\lambda_j T}, \quad m = |w|.$$

We claim this as a **bound on the form**, not a tight inequality: the constants c_j are not estimated, and we instead **measure the consequence** — that channel j is certified only to horizon

$$T_j(\epsilon) \sim \frac{1}{\lambda_j} \log \frac{1}{\epsilon} \quad (\lambda_j > 0), \quad T_j = \infty \quad (\lambda_j \leq 0),$$

which §5.2 recovers to within 0.4%. As $\epsilon_{\max}, \delta \rightarrow 0$ (exact equivariance) the configuration term vanishes and Theorem A is recovered.

Corollary (the certified region). The set $\mathcal{C} = \{(w, T, \epsilon) : |\Delta \text{Err}| \leq \epsilon\}$ is the **coarse** (ϵ large), **slow** ($\lambda_j \leq 0$), **low- $|w|$ corner**; its boundary is set by $\langle S \rangle$ (configuration) and the spectrum $\{\lambda_j\}$ (horizon and resolution). With *exact* equivariance the configuration axis is unbounded and only the spectrum limits horizon $\times$ resolution — a trade-off horizon $\times$ demanded-resolution $\lesssim \text{const}(\{\lambda_j\})$. This is the formal content of “ultra-long forecasts must be coarse”: eclipses for millennia (conserved $\lambda \leq 0$ actions), weather at $\sim$ two weeks (chaotic $\lambda > 0$).¹

3.3 Scale versus structure, quantified

The tagline *scale buys interpolation; structure buys a certificate* can be made precise as a statement about the **certified region** — the inputs a learner can guarantee, in the worst case over targets consistent with what it observed on a training set T . Formally $\mathcal{C} = \{z : \inf_f \sup_{\Phi} \|f(z) - \Phi(z)\| \leq \epsilon\}$, where f ranges over the learner’s hypotheses (exact on T) and Φ over the admissible targets agreeing on T .

Proposition (separation). (*Structure.*) Under the equivariant prior ($\Phi \in \mathcal{D}_{G,f}$ equivariant and exact on T , S generating G), Theorem A makes the error orbit-constant, so $\mathcal{C} \supseteq G \cdot T$: the **entire orbit is certified to error 0, independent of ϵ** , from the $k = |S|$ generator checks. (*Scale/data.*) Under the equivariance-free prior of L -Lipschitz targets consistent on T , the McShane extensions $\Phi_{\pm}(z) = \min / \max_{t \in T} (\Phi(t) \pm L \text{dist}(z, t))$ are admissible and differ at z by up to $2L \text{dist}(z, T)$, so every learner has minimax error $\geq L \text{dist}(z, T)$ and $\mathcal{C} \subseteq \{z : \text{dist}(z, T) \leq \epsilon/L\}$ — an ϵ/L -tube around the training set that **collapses to T as $\epsilon \rightarrow 0$** .

The separation is sharp in its dependence on ϵ : a point at orbit-distance D from T — the unseen far end of the wedge in Experiments 7 and 9 — is certified by *structure* for **every** $\epsilon > 0$, but by *scale/data* only once $\epsilon \geq LD$. Scaling a non-equivariant model (more capacity, more in-wedge data) sharpens its fit *inside* the tube; it cannot enlarge the tube, because the bound $L \text{dist}(z, T)$ is information-theoretic (independent of model size). This is the precise content of the empirical curves: the equivariant model is flat over the whole orbit while every baseline scale climbs once $\text{dist}(z, T)$ exceeds ϵ/L .

4. The Noether Hinge

Theorems A and B hold for *whatever* channels happen to be slow. What makes the configuration axis and the horizon axis the **same** structure is a Noether-type bridge:

Conjecture (group $\Rightarrow$ slow). When the symmetry is a symmetry of the **dynamics** (not merely the encoder), the group-invariant ($\ell=0$) channels coincide with the conserved/slow ($\lambda_j \leq 0$) channels. Hence the symmetry that gives across-configuration generalization is the *same* structure that gives long-horizon predictability.

This is **not automatic** — invariant $\neq$ conserved in general — so we treat it as a *falsifiable, measured* conjecture. The defensible direction is the inclusion **slow** $\subseteq$ **invariant** $\oplus$ **conserved-equivariant**, not the converse: the conserved (slowest) modes live in the invariant-or-conserved subspace, but not every invariant channel is slow (a rotation-invariant $|r|^2$ oscillates). The mechanism is Noether’s. A continuous symmetry of the dynamics yields a conserved current; a

¹An interpretive aside, secondary to the spectral law: a long-range text that asserts only *regime-level* change rather than dated specifics is, structurally, making the *correct* move for a high- λ system, where only the coarse/invariant component is certifiable.

conserved *scalar* charge ($\ell=0$ Casimir, e.g. energy) **is** group-invariant, while a conserved *non-scalar* charge (e.g. three-dimensional angular momentum L , an $\ell=1$ vector) is *equivariant*, **not** invariant. The slow subspace therefore lies in invariant $\oplus$ conserved-equivariant, which **collapses to the invariant component exactly when every conserved charge is a scalar** — the two-dimensional case, where $\text{slow} \subseteq \text{invariant}$ holds literally. Section 5.3 measures both forms: the clean 2D containment, and its 3D refinement in which angular momentum is recovered from the equivariant block at polynomial degree two (the cross product).

5. Experiments

All experiments are CPU/1-GPU-scale, run with explicit random seeds, and **gated**: a run reports INCONCLUSIVE rather than loosen a threshold. Load-bearing results are additionally guarded by unit tests. The experiment-to-code map, seed counts, and headline numbers are collected in Appendix A; multi-seed ranges below are seed min–max over three seeds unless noted.

5.1 The configuration axis

Experiment 1 (exact compositional flatness). On a structured object-interaction world with $\text{SE}(3) \times S_n$ symmetry, the equivariant model’s whole-pipeline relMSE is **exactly** $\times 1.00$ **flat over composition words** $m = 0, \dots, 8$ (constant 0.2625), while a matched non-equivariant baseline blows up by $\times 12.89$. The predictor’s Jacobian spectrum is measured directly: **48 of 96 channels are contractive** ($\sigma \leq 1$, certified long-horizon) versus 48 expansive (max $\sigma = 49$, min $\sigma = 0.003$) — the clean coarse/fine split that is exactly Theorem B’s input. A unit test verifies Theorem A holds *architecturally* at initialization (word-deviation 1.2×10^{-7} , before any training) while the control deviates 20%.

Experiment 2 (exponential certification on $\mathbb{Z}_2^6$). The 64 hexagrams of the *I Ching* are exactly $\{\pm 1\}^6 = \mathbb{Z}_2^6$, with the 6 changing-lines as a generating set. Training a sign-equivariant per-line predictor $f(z, a)_i = h_\theta(a_i) \odot z_i$ on **only the 6 single-line generators** (plus the identity, 7 of 64 actions) certifies it over **all $2^6=64$ compositions** to machine precision (worst relMSE $\sim 10^{-33}$; equivariance residual exactly 0), while a non-equivariant baseline trained on the same 7 actions degrades monotonically with composition length ($1.6\text{--}1.7 \times 10^{-5}$ on seen actions $\rightarrow 0.59$ at six flips, Figure 1). This makes “ k generator checks $\Rightarrow 2^k$ certified set” literal, with genuine learning and an honest contrast.

5.2 The horizon $\times$ resolution staircase

Experiment 3 (the predictability staircase). A learned one-step predictor on a designed multi-Lyapunov system — a conserved invariant ($\lambda = -\infty$), a contracting channel ($\lambda = \ln 0.75$), a neutral $\text{SO}(2)$ rotor ($\lambda = 0$), and a chaotic angle-doubling channel $z \mapsto z^2$ ($\lambda = \ln 2$) — is rolled out autoregressively. The model **recovers the chaotic Lyapunov exponent to within 0.4%** ($\hat{\lambda} = 0.690\text{--}0.692$ across three seeds, versus $\ln 2 = 0.693$), and its per-channel certified horizon obeys $T_j(\epsilon) \sim \log(1/\epsilon)/\lambda_j$: the chaotic “fine detail” is certifiable for only 3–10 steps over the full ϵ grid, and its horizon grows *only logarithmically* as ϵ coarsens (measured slope $dT/d \log \epsilon = 1.3\text{--}1.6 \approx 1/\lambda = 1.44$), while the **contracting channel stays certified for the entire 90-step rollout at every ϵ** (the conserved and rotor channels too, except at the finest $\epsilon=0.005$, where they fall to $\sim 40\text{--}90$ steps). At $\epsilon=0.05$ the slow channels outlast the chaotic detail by 12–15 $\times$ (Figure 2). This is the predictability staircase: long horizon $\Rightarrow$ coarse and invariant only.

5.3 The Noether hinge

Experiment 4 (the hinge in 2D). On a two-dimensional $\text{SO}(2)$ central-force system — conserved energy E and angular momentum L are invariant scalars, the orbital phase is fast — a learned equivariant autoencoder-with-latent-dynamics (a hand-built 2D Vector-Neuron model) spontaneously organizes the conserved quantities into its architecturally-invariant ($\ell=0$) block. Regression R^2 for (E, L) is 0.92–0.99 **from the invariant block versus ≤ 0.012 from the equivariant ($\ell=1$) block**, and the slowest mode the invariant block admits (0.009–0.031 across seeds) is 4.7–17 $\times$ slower than anything the equivariant block admits (0.145, exactly the orbital phase velocity $2 \sin(\omega \Delta t/2)$). We read the containment off this min-mode gap — the equivariant block’s *slowest* achievable mode is already fast, so no slow direction escapes the invariant block in this system — rather than by a direct subspace test. The payoff is the certificate: the equivariant model’s slow subspace **is** its invariant subspace, whose out-of-distribution group-action residual is $\sim 10^{-16}$, exact and architectural for all of $\text{SO}(2)$. A matched non-equivariant baseline also learns E, L ($R^2 = 0.99$) but smears them




Figure 1: Training on the 6 generators of $\mathbb{Z}_2^6$ certifies all 64 compositions (equivariant, machine precision) while a non-equivariant baseline degrades with composition length.




Figure 2: Left: rollout error growth recovers the Lyapunov spectrum. Right: the certified-horizon staircase $T_j(\epsilon) \sim \log(1/\epsilon)/\lambda_j$.

across directions that **drift by** ≈ 1.17 under the same group action, and carries no certificate (Figure 3). The honest non-converse is visible in the data (the invariant but fast scalar $|r|^2$): $\text{slow} \subseteq \text{invariant}$, not $=$.




Figure 3: Left: the slowest latent modes live in the invariant ($\ell=0$) block. Right: the certificate — the invariant subspace stays group-invariant to 10^{-16} where the non-equivariant model’s slow directions drift by ~ 1 .

Experiment 5 (lift to a 3D contact interaction). Pushing the hinge to a more embodied regime — two bodies in a three-dimensional well with a **soft pairwise repulsion (contact)**, modeled by an $\text{SO}(3)$ -equivariant $\ell=0 \oplus \ell=1$ Vector-Neuron — gives a *partial* lift, honestly reported. The *Noether-content* half lifts: the invariant $\ell=0$ block recovers the conserved (E , contact-distance) at $R^2=0.60\text{--}0.86$ versus ≤ 0.06 for the $\ell=1$ block, even with contact. But the clean containment $\text{slow} \subseteq \text{invariant}$ does **not** lift (invariant slowest mode $0.024 \approx 0.026$ for the equivariant block; this gate is reported INCONCLUSIVE) — for a principled reason: in 3D, angular momentum $L = \sum_i r_i \times v_i$ is a conserved $\ell=1$ *vector*, so the equivariant block too carries a slow conserved mode. The clean 2D containment (where L is a scalar) becomes, in 3D, $\text{slow} \subseteq (\text{invariant} \oplus \text{conserved-equivariant})$. The Noether-content claim is dimension-robust; the clean-containment form is 2D-specific — a sharpening of scope, not a failure.

Experiment 6 (3D-aware containment). The conserved physics splits by isotypic *type* and polynomial *degree*: E (an invariant quadratic) is recovered linearly from the $\ell=0$ block ($R^2=0.62\text{--}0.91$ across seeds), whereas $L = \sum_i r_i \times v_i$ (a conserved $\ell=1$ vector) is **bilinear** — not linear in either block (≤ 0.04) but recovered at $R^2=1.00$ (range $0.998\text{--}1.000$, the robust load-bearing result) by the **degree-two cross-product** readout of the $\ell=1$ block. So $\text{slow} \subseteq (\text{invariant} \oplus \text{conserved-equivariant})$ holds exactly in 3D, with the equivariant conserved part accessed at degree two — and $r \times v$ is exactly the cross product a degree-one Vector-Neuron cannot form, tying the hinge’s 3D form to the companion paper’s bilinear-message result. The full gate, which additionally demands E -recovery $R^2>0.8$ in the $\ell=0$ block, passes on one of three seeds — as honestly surfaced as the INCONCLUSIVE gate of Experiment 5 — while the degree-two L recovery is robust across all three.

5.4 Structure versus scale

Experiment 7 (structure versus scale). Train on a 50° wedge of an $\text{SO}(2)$ orbit and test around the full circle. The exactly-equivariant model’s error is **flat over the entire orbit** (out-of-wedge / in-wedge ratio $1.1\text{--}1.2$) — a certificate that holds from partial data by the identity $\text{err}(R_\theta s) = \text{err}(s)$. A non-equivariant baseline **scaled across an $88\times$ parameter range** ($3.8\text{k} \rightarrow 337\text{k}$) buys genuine *interpolation*: its in-wedge error drops $31\text{--}166\times$ and even **beats** the $56\times$ -smaller equivariant model in-distribution — but out of the wedge the largest baseline remains $10\text{--}155\times$ worse than the equivariant model and never reaches its floor (Figure 4). The fair nuance — *scale can win in-distribution* — is exactly why the contribution is the **certificate** (an orbit-wide guarantee from partial data), not a per-point accuracy contest.




Figure 4: Left: error over the SO(2) orbit (equivariant flat; the scaled baselines dip below it in-wedge then climb out). Right: out-of-wedge error versus baseline scale plateaus far above the equivariant floor.

5.5 Approximate symmetry

Experiment 8 (graceful degradation and a measured threshold). We break the *world's* SO(2) symmetry with an anisotropy knob β (the potential becomes $V = \frac{1}{2}(x^2 + (1 + 2\beta)y^2)$, so angular momentum is no longer conserved) and re-run the wedge-train / full-circle-test. Three findings, robust across seeds: (i) at $\beta=0$ the certificate is exact — the equivariant model is 68–320× **better out-of-wedge** than the baseline; (ii) as the measured world symmetry-defect ϵ_{world} grows, the equivariant out-of-wedge error grows **smoothly and monotonically** (correlation 0.88–0.98, exactly Theorem B's ϵ_{max} term — a graceful slope, not a cliff); and (iii) the equivariant model keeps beating the non-equivariant one up to a **measured symmetry-content threshold** ($\epsilon_{\text{world}} \approx 0.01$ –0.06, seed-dependent), beyond which the now-wrong symmetry assumption hurts more than it helps (Figure 5). So *approximate* symmetry buys an *approximate* certificate with exactly the error budget the theory predicts, and the boundary where structure stops paying is itself measured, not assumed.




Figure 5: Left: equivariant out-of-wedge error grows $\propto \epsilon_{\text{world}}$ (Theorem B's ϵ term). Right: the equivariant model beats the baseline out-of-wedge up to a symmetry-content threshold, then crosses over.

5.6 From certificate to action

A certificate is most useful if one need neither know the group in advance nor a decoder to act on it. Two results, re-framed from a companion line of experiments (not fresh runs here), close that loop. **Discovery:** from a blank slate of K learnable 3×3 matrices — nothing antisymmetric or bracket-closing imposed — gradient descent *rediscovers* a frozen teacher’s symmetry algebra ($\mathfrak{so}(3)$, dimension 3, on a rotationally-symmetric teacher; the smaller $\mathfrak{so}(2)_z$, dimension 1, on a rotation-broken one — it reports the *surviving* group and refuses to invent the rest), and distilling those discovered generators into a free predictor buys back most of the across-group generalization the hard-wired equivariant model gets for nothing (closing more than half the out-of-distribution gap, matching the hand-wired oracle, transferring *exactly* the symmetry discovered). So the generating set S that anchors the configuration axis need not be postulated — it can be *measured*, falsifiably. **Generation:** given S , one can traverse the certified orbit to an unseen composed goal and execute it **closed-loop without a decoder** — latent-goal reaching driven purely by the equivariance identity over 24 paired seen-vs-unseen $SE(3)$ -orbit tasks, closing ~ 0.59 of the orientation gap on the best deployable variant and — the load-bearing point — achieving the *same* success on the unseen orbit as on the seen one (unseen/seen ratio 1.000): *exact transfer*, not high absolute success. Together: *discover* $S \rightarrow$ *certify* $\langle S \rangle \rightarrow$ *generate* $\rightarrow$ *act* — the certificate is not merely descriptive but a usable plan-and-execute criterion.

5.7 Beyond toys: the certificate on real contact dynamics

Experiment 9 (PushT, real physics-engine contact dynamics). Every experiment so far uses dynamics we *constructed* to be equivariant — the standing concern that the symmetry is built into a toy. We close that gap on **PushT**, a planar pushing benchmark whose contact dynamics are produced by a 2D physics engine we did not author. In the interior (the block kept away from the workspace walls) the dynamics are genuinely $SO(2)$ -equivariant — rotating the whole scene rotates the physics. On a $\pm 50^\circ$ wedge of real interior transitions we train an $SO(2)$ -equivariant world model — an invariant-scalar-gated Vector-Neuron, exactly equivariant to residual $\sim 10^{-7}$, so its contact features (agent–block distance and contact angle) live in the invariant pathway — and, as a *fair* baseline, non-equivariant MLPs across a $160\times$ parameter ladder ($1.7k \rightarrow 272k$); both are trained one-step with identical data, optimiser, and cosine schedule. We then roll each model out $H=10$ steps and measure rollout relMSE across the full orbit of scene orientations (rotating a held-out real test set, as in Experiment 7).

Three findings, robust across 3 seeds (Figure 6): (i) the equivariant model’s rollout error is **exactly flat over the orbit** (out-of-wedge / in-wedge ratio 1.00 at every horizon) — the certificate, now on a *learned* model of *real* contact physics; (ii) it is **competitive in-distribution** (in-wedge relMSE 0.13–0.15 vs the best MLP’s 0.14–0.19), so the certificate is not bought by being a worse predictor — “flat” is also “good” here; and (iii) **no MLP scale reaches the equivariant floor out-of-wedge** — even the 272k-parameter MLP ($16\times$ the equivariant model’s size) stays $2.1\text{--}3.9\times$ above it. The gap is smaller than the toy’s because a single PushT step is easy to predict in-distribution; the certificate’s value is precisely that it survives the *rollout*, out of distribution, where scale does not follow. This is the keystone non-toy result: the configuration certificate holds for a learned model of dynamics we did not design.

6. Related Work

Equivariance and constant error (our companion line). A companion paper establishes, for a single group element and one resolution, that exact equivariance makes one-step error constant across the group, with exact-flatness and closed-loop-invariance corollaries. The present paper lifts that corner along three axes: Theorem A is its multi-step, full-monoid generalization; Theorem B adds the horizon $\times$ resolution stratification; and Lemma 1 turns k generator checks into an exponential certified set.

Equivariant predictors. Block-rotation predictors (BRo-JEPA, arXiv:2606.01372) match a $\mathbb{Z}/10\mathbb{Z}$ cyclic structure to the latent — “add k ” becomes “rotate $k\theta$ ” — and report 99.46% best-config zero-shot transfer on MNIST modular arithmetic; unitary-predictor world models (UWM-JEPA, arXiv:2605.25313) do the analogous thing with $U(d)$. Theorem A explains *why* such zero-shot transfer occurs (the representation cancels inside the norm — orthogonal for the former, unitary for the latter), and the closure lemma quantifies *how far* it reaches (the whole generated monoid from its generators). These works share our toy-construction caveat, and we cite them as cross-group mechanism corroboration rather than scale results.




Figure 6: Experiment 9 (PushT, real contact dynamics). Left: 10-step rollout relMSE over the orbit of scene orientations — the learned SO(2)-equivariant model is exactly flat (ratio 1.00) and competitive in-distribution, while non-equivariant baselines dip in-wedge then climb out. Right: across a 160× parameter ladder no baseline reaches the equivariant floor out-of-wedge, at any rollout horizon.

Latent-geometry regularizers. LeJEPa (Balestriero and LeCun, 2025) pushes the latent toward an isotropic Gaussian; UR-JEPa (Le, 2026; arXiv:2606.01443) challenges that target, arguing isotropy conflicts with the manifold hypothesis and that a data-discovered anisotropy is preferable. Our latent’s anisotropy is neither isotropic nor data-discovered but **group-prescribed**: the covariance is pinned to the representation, $\rho \Sigma \rho^\top = \Sigma$ (measured residual 3×10^{-4} , versus 1.04 for a non-equivariant baseline). In a head-to-head (companion line, single seed) the data-fit anisotropy is best *in-distribution*, but only the group-prescribed anisotropy transfers out of distribution (a $\sim 320\times$ gap, consistent with the three-seed 68–320× and 10–155× of Experiments 8 and 7). The certificate is therefore complementary — a first-order, per-situation guarantee from the group, not a second-order distributional prior — and it predicts precisely *when* a discovered anisotropy will fail to generalize (off the training orbit).

Oracle-bypass diagnostics. IMWM (Gao et al., 2026; arXiv:2606.01626) uses the same falsifiable move as our capacity-ladder experiments — replace the learned model with oracle dynamics and see what is still missing — and finds a finite-budget planner *still* fails. It attributes the residual to **search** (proposal-sampling volume), whereas our analogous diagnosis attributes it to **representation** (permutation-invariant encoder pooling). These are complementary bottlenecks: a predictability certificate bounds the *model’s* error over $\langle S \rangle \times T \times \epsilon$, and makes no claim about the planner’s search budget downstream.

Geometry on the action side. LDA (“The Lie We Tell”, Chuang et al., 2026; arXiv:2606.01847) names the *Euclidean Fallacy* — flattening an SE(3) pose into $\mathbb{R}^{12}$ breaks the manifold constraint, the coordinate-change equivariance, and geodesic optimality — and fixes it by running diffusion *on* SE(3) (tangent-space score, exponential-map retraction). It states our core motivation — *do not flatten geometric quantities* — on the action side. As a diffusion policy its gains are modest but robust (CALVIN average task length $3.27 \rightarrow 3.51$, +7.3%), the same profile as equivariant methods generally; our framing explains why a geometric prior buys a *kind* of guarantee (consistency across the group, robustness out of distribution) rather than a uniform accuracy jump.

Predictability horizons. The $T(\epsilon) \sim \log(1/\epsilon)/\lambda$ law is classical for dynamical systems (Lyapunov; numerical weather prediction). Our contribution is to measure it on a *learned latent world model*, tie its slow subspace to the group-invariant subspace via the Noether hinge, and fold it into a single multi-axis certificate for equivariant models.

7. Limitations

- **Scale and scope.** All experiments are CPU/1-GPU. Experiment 9 validates the central claim on a real physics-engine contact simulator (PushT) — dynamics we did not author — but still on *structured state* (not pixels), a single $SO(2)$ task, and at a *modest* out-of-distribution gap (2–4 $\times$, large only because a single PushT step is easy to predict in-distribution). We claim a mechanism and a tool, demonstrated cleanly, not a scaled benchmark.
- **Scope of the exact certificate.** Theorem A requires (A3): the group must be a symmetry of the *dynamics*, not merely the encoder. The exact certificate therefore holds where the group is a genuine dynamical symmetry (orbital and conservative systems, free space, idealized manipulation). Everywhere else one is in Theorem B’s approximate regime, whose degradation we *measure* rather than assume: Experiment 8 shows it is graceful ($\propto \epsilon_{\text{world}}$) up to a measured threshold $\epsilon_{\text{world}} \approx 0.01\text{--}0.06$ (seed-dependent; one seed crosses over as early as $\epsilon \approx 0.008$). The honest reading is “exact on symmetric dynamics; gracefully approximate, with a measured boundary, elsewhere.”
- **Theorem B is a bound on the form, not a tightness proof.** The constants c_j are not estimated, and the *isotypic* refinement (channels chosen by ℓ -type, which ties the spectrum to equivariance and the hinge) is asserted, not separately measured — Experiment 3 measures the scalar law and Experiment 1 splits contractive from expansive but not by ℓ -type.
- **The Noether hinge is a measured conjecture**, confirmed on controlled systems with the honest non-converse (invariant $\nRightarrow$ slow). Its lift to a *real, contact-rich embodied* model — where symmetry is approximate and the latent is learned end-to-end — is the primary open problem; the 3D contact experiment is a clean two-body toy whose clean-containment gate is INCONCLUSIVE (resolved in type-and-degree form by Experiment 6).
- **§5.6 re-frames companion-line results**, not fresh runs: the decoder-free reach is *exact transfer* (unseen/seen ratio 1.000) at a *partial* absolute success (~ 0.59 of the orientation gap). We report the transfer ratio as the load-bearing claim, not absolute task success.
- **“Flat” is not “good.”** A constant error across the group says nothing about its magnitude; Theorem A certifies *consistency*, and the “scale never reaches the certificate” claim is an *architectural* fact, **proved** in Lemma 2 (orbit-constant error against every equivariant target $\iff$ equivariance) and quantified in §3.3 (a non-equivariant L -Lipschitz learner certifies only an ϵ/L -tube around the data) — not the conclusion of a finite sweep.

Falsifiability. The framing makes refutable predictions, not after-the-fact rationalizations. A finite non-equivariant network attaining exact orbit-flatness would break Theorem A’s architectural premise; a symmetric-dynamics system whose conserved/slow modes lay *outside* the invariant $\oplus$ conserved-equivariant subspace would break the hinge; and a channel whose certified horizon failed to scale as $\log(1/\epsilon)/\lambda_j$ would break Theorem B. We have not observed these; they are the experiments that would sink the thesis.

8. Conclusion

For equivariant world models, structure delivers a *kind* of result scaling cannot: a **predictability certificate** — a provable, computable, training-free region across configuration, horizon, and resolution. We established the master theorem (Theorem A, the exact configuration certificate, proved — its closed-loop clause under the assumed equivariant-planner condition; Lemma 2, its converse, characterizing the certificate *as* equivariance; Theorem B, the spectral horizon $\times$ resolution law, stated as a bound and measured), exhibited the certified region as the coarse-invariant-slow-low-composition corner, measured the Noether hinge that unifies the symmetry and time axes, and showed — empirically and via the quantitative separation of §3.3 — that an 88 $\times$ -scaled non-equivariant model buys interpolation but never the certificate, an architectural impossibility rather than an unfinished sweep. *Scale buys interpolation; structure buys a certificate.* The same criterion that tells you which compositions a robot policy will handle zero-shot also tells you why eclipses are forecastable for millennia while weather is not — one structural law, read across domains.

References

Concurrent and prior work referenced above (arXiv identifiers from the project’s source notes; external numbers are quoted from the cited works):

- **BRO-JEPA** — block-rotation $\mathbb{Z}/10\mathbb{Z}$ predictor; 99.46% best-config zero-shot on MNIST modular arithmetic. arXiv:2606.01372.
- **UWM-JEPA** — unitary-predictor ($U(d)$) belief-space world model. arXiv:2605.25313.
- **UR-JEPA** (Le, 2026) — uniform-rectifiability latent regularizer; manifold-hypothesis critique of isotropic SIGReg. arXiv:2606.01443.
- **IMWM** (Gao et al., 2026) — oracle-bypass diagnosis localizing the residual to planner *search*. arXiv:2606.01626.
- **LDA / “The Lie We Tell”** (Chuang et al., 2026) — the *Euclidean Fallacy*; SE(3)-intrinsic diffusion policy; CALVIN average task length $3.27 \rightarrow 3.51$ (+7.3%). arXiv:2606.01847.
- **LeJEPA** (Balestrierio and LeCun, 2025) — SIGReg / isotropic-Gaussian self-supervised objective (the latent-geometry target that UR-JEPA, and our group-prescribed anisotropy, depart from).

The $T(\epsilon) \sim \log(1/\epsilon)/\lambda$ predictability-horizon law is classical (Lyapunov / Lorenz; standard numerical-weather-prediction practice).

Appendix A. Reproducibility

Every experiment sets random seeds explicitly, prints an INCONCLUSIVE verdict rather than loosen a gate, and writes its figure and JSON to `papers/figures/`. The multi-seed steps commit per-seed JSONs (`papers/figures/step5*_seeds.json` and `step59_pusht_certificate_seeds.json`, seeds 0/1/2), regenerated by `experiments/aggregate_seeds.py`; every range quoted above is the seed min–max from those files. The configuration-flatness experiment (Experiment 1) self-aggregates its three seeds into the means in `step47_certificate.json`. The full test suite passes together; `tests/conftest.py` isolates the float64 experiments from the float32 codebase.

Exp.	Result	Code	Test	Seeds	Headline number
1	Exact compositional flatness + spectrum	<code>experiments/step47_certificate.py</code>	<code>tests/test_step47_certificate.py</code>	3	$\text{relMSE} \times 1.00$ flat ($m=0..8$); 48/96 contractive
2	$\mathbb{Z}_2^6$ exponential certificate	<code>experiments/step49_iching_certificate.py</code>	<code>tests/test_step49_iching_certificate.py</code>	1	6 generators $\rightarrow$ all 64; worst $\text{relMSE} \sim 10^{-33}$
3	Horizon $\times$ resolution staircase	<code>experiments/step52_horizon_resolution.py</code>	<code>tests/test_step52_horizon_resolution.py</code>	3	$\hat{\lambda}=0.69$ vs $\ln 2$; slope $\approx 1/\lambda$
4	Noether hinge (2D)	<code>experiments/step50_noether_hinge.py</code>	<code>tests/test_step50_noether_hinge.py</code>	3	R^2 0.92–0.99 vs ≤ 0.012 ; cert 10^{-16} vs 1.17
5	Lift to 3D contact	<code>experiments/step57_embodied_hinge.py</code>	<code>tests/test_step57_embodied_hinge.py</code>	3	clean- containment gate INCONCLUSIVE ; Noether content lifts (invariant $R^2=0.60$ –0.86 vs ≤ 0.06); resolved by Exp. 6

Exp.	Result	Code	Test	Seeds	Headline number
6	3D-aware containment	experiments/step58_3d_containment.py		3	$E \rightarrow \ell=0$ linear ($R^2=0.62\text{--}0.91$); L bilinear $\rightarrow \ell=1$ degree-2 cross ($R^2=1.00$, range 0.998–1.000)
7	Structure vs scale	experiments/step51_structure_vs_scale.py		3	equiv flat 1.1–1.2; in-wedge gain 31–166 $\times$; best baseline 10–155 $\times$ above floor
8	Approximate symmetry	experiments/step53_approximate_symmetry.py		3	out-of-wedge cert exact at $\beta=0$ (68–320 $\times$); graceful $\propto \epsilon$ (corr 0.88–0.98); threshold $\epsilon \approx 0.01\text{--}0.06$
9	Certificate on real contact dynamics (PushT)	experiments/step59_pusht_certificate.py		3	learned equivariant exactly flat over the orbit (ratio 1.00, equiv-resid $\sim 10^{-7}$) and competitive in-dist; no MLP scale (1.7k–272k) reaches the floor out-of-wedge (2.1–3.9 $\times$, $H=10$)